

Rust Programming Language

Yumcoder

University of Toronto

December 29, 2022

Contents

- 1 Introduction
- 2 Common Programming Concepts
- 3 Understanding Ownership
- 4 Using Structs to Structure Related Data
- 5 Enums and Pattern Matching
- 6 Managing Projects with Packages, Crates, and Modules
- 7 Common Collections
- 8 Error Handling
- 9 Generic Types, Traits, and Lifetimes
- 10 Writing Automated Tests
- 11 An I/O Project: Building a Command Line Program
- 12 Functional Language Features: Iterators and Closures
- 13 More About Cargo and Crates.io
- 14 Smart Pointers
- 15 Fearless Concurrency
- 16 Object-Oriented Programming Features of Rust
- 17 Patterns and Matching
- 18 Advanced Features
- 19 Multithread Web Server
- 20 Tokio

Introduction

Installing rustup on Linux or macOS

If you're using Linux or macOS, open a terminal and enter the following command:

```
$ curl --proto '=https' --tlsv1.3 https://sh.rustup.rs -sSf | sh
```

If the install is successful, the following line will appear:

```
Rust is installed now. Great!
```

To check whether you have Rust installed correctly, open a shell and enter this line:

```
$ rustc --version
```

Updating, Uninstalling and Local Documentation

Once Rust is installed via rustup, when a new version of Rust is released, updating to the latest version is easy. From your shell, run the following update script:

```
$ rustup update
```

To uninstall Rust and rustup, run the following uninstall script from your shell:

```
$ rustup self uninstall
```

The installation of Rust also includes a local copy of the documentation, so you can read it offline. Run `rustup doc` to open the local documentation in your browser.

Hello, World!

Open a terminal and enter the following commands:

```
1 $ mkdir ~/projects
2 $ cd ~/projects
3 $ mkdir hello_world
4 $ cd hello_world
```

Make a new source file and save it as main.rs:

```
1 fn main() {
2     println!("Hello, world!");
3 }
```

Compile and run the file:

```
$ rustc main.rs
$ ./main
Hello, world!
```

Cargo

- Cargo is **Rust's build system and package manager**.
- Most **Rustaceans** use this tool to manage their Rust projects because Cargo handles a lot of tasks for you, such as building your code, downloading the libraries your code depends on, and building those libraries.
 - 👉 **Rustaceans** are people who use Rust, contribute to Rust, or are interested in the development of Rust.
- Cargo comes installed with Rust if you used the official installer.

```
$ cargo --version
```

Hello, Cargo!

Creating a Project with Cargo:

```
$ cargo new hello_cargo  
$ cd hello_cargo
```

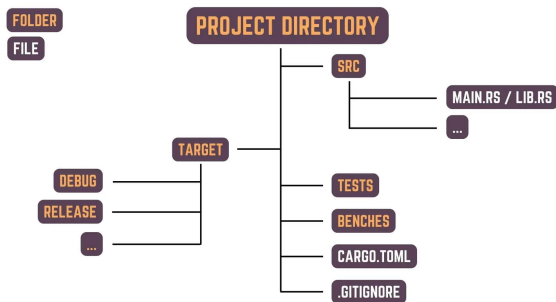


Figure: Folder structure for projects in rust programming language.

Cargo.toml

This file is in the **TOML** (Tom's Obvious, Minimal Language) format, which is **Cargo's configuration** format.

```
1 [package]
2 name = "hello_cargo"
3 version = "0.1.0"
4 edition = "2021"
5
6 # See more keys and their definitions at
6 ↪ https://doc.rust-lang.org/cargo/reference/manifest.html
7
8 [dependencies]
```

Building and Running a Cargo Project

Now open `src/main.rs` and take a look:

```
1 fn main() {  
2     println!("Hello, world!");  
3 }
```

Build your project by entering the following command:

```
$ cargo build  
Compiling hello_cargo v0.1.0 (file:///projects/hello_cargo)  
Finished dev [unoptimized + debuginfo] target(s) in 2.85 secs
```

- Use `cargo run` to compile and then run (all in one command).
- When your project is finally ready for release, you can use `cargo build --release` to compile it with optimizations (create an executable in `target/release` instead of `target/debug`).

Common Programming Concepts

Variables and Mutability

- by default variables are **immutable**.
- This is one of many nudges **Rust** gives you to write your code in a way that takes advantage of the **safety and easy concurrency** that Rust offers.
 - However, you still have the option to make your variables **mutable**.

Immutable variables

When a variable is immutable, *once a value is bound to a name, you can't change that value.*

```
1 fn main() {  
2     let x = 5;  
3     println!("The value of x is: {x}");  
4     x = 6;  
5     println!("The value of x is: {x}");  
6 }
```

Immutable variables (2)

You should receive an error message, as shown in this output:

```
$ cargo run
   Compiling variables v0.1.0 (file:///projects/variables)
error[E0384]: cannot assign twice to immutable variable `x`
  --> src/main.rs:4:5
   |
2 |     let x = 5;
   |     -
   |     |
   |     first assignment to `x`
   |     help: consider making this binding mutable: `mut x`
3 |     println!("The value of x is: {x}");
4 |     x = 6;
   |     ^^^^^ cannot assign twice to immutable variable
```

For more information about this error, try `rustc --explain E0384`.
error: could not compile `variables` due to previous error

Mutable variables

Although variables are immutable by default, you can make them mutable by adding `mut` in front of the variable name.

```
1 fn main() {  
2     let mut x = 5;  
3     println!("The value of x is: {x}");  
4     x = 6;  
5     println!("The value of x is: {x}");  
6 }
```

Mutable variables (2)

```
$ cargo run
  Compiling variables v0.1.0 (file:///projects/variables)
    Finished dev [unoptimized + debuginfo] target(s) in 0.30s
    Running `target/debug/variables`
The value of x is: 5
The value of x is: 6
```


Constants

Rust's naming convention for constants is to use *all uppercase with underscores between words*.

```
const THREE_HOURS_IN_SECONDS: u32 = 60 * 60 * 3;
```

Like immutable variables, constants are values that are bound to a name and are not allowed to change, but there are a few differences between constants and variables.

- you aren't allowed to use `mut` with constants. Constants aren't just immutable by default—**they're always immutable**. You declare constants using the `const` keyword instead of the `let` keyword, and **the type of the value must be annotated**.
- constants may be set only to a constant expression, **not the result of a value that could only be computed at runtime**.

Shadowings

Rustaceans say that the first variable is shadowed by the second, which means that the second variable is what the compiler will see when you use the name of the variable. We can shadow a variable by using the same variable's name and repeating the use of the `let` keyword as follows

```
1 fn main() {  
2     let x = 5;  
3  
4     let x = x + 1;  
5  
6     {  
7         let x = x * 2;  
8         println!("The value of x in the inner scope is: {x}");  
9     }  
10  
11     println!("The value of x is: {x}");  
12 }
```

Shadowing (2)

```
$ cargo run
  Compiling variables v0.1.0 (file:///projects/variables)
    Finished dev [unoptimized + debuginfo] target(s) in 0.31s
    Running `target/debug/variables`
The value of x in the inner scope is: 12
The value of x is: 6
```

Shadowing vs Mutable variables

We're effectively creating a new variable when we use the `let` keyword again, we can change the type of the value but reuse the same name.

```
1 fn main() {  
2     let spaces = "    ";  
3     let spaces = spaces.len();  
4 }
```

The first `spaces` variable is a string type and the second `spaces` variable is a number type.

```
1 fn main() {  
2     let mut spaces = "    ";  
3     spaces = spaces.len();  
4 }
```

we'll get a **compile-time error**:
mismatched types, expected `'&str'`,
found `'usize'`.

Data Types - Primitive type

Data type subsets:

- Scalar
 - Integers
 - Floating-point numbers
 - Booleans
 - Characters
- Compound
 - Tuples
 - Arrays

```

1 let y: f32 = 3.0;
2 // y is variable name
3 // f32 is IEEE-754 double
  ↳ precision standard variable
  ↳ type
4 // 3.0 is initial value
  
```

Statically typed language

Keep in mind that Rust is a statically typed language, which means that **it must know the types of all variables at compile time.**

The compiler can usually **infer** what type we want to use based on the value and how we use it (example: `let x = 2.0; // f64`).

Integer Types

- An integer is a number without a fractional component.
- This type declaration indicates that the value it's associated with should be an unsigned integer (signed integer types start with **i**, instead of **u**) that takes up 32 bits of space.

Length	Signed	Unsigned
8-bit	i8	u8
16-bit	i16	u16
32-bit	i32	u32
64-bit	i64	u64
128-bit	i128	u128
arch ¹	isize	usize

¹the isize and usize types depend on the architecture: 64 bits if you're on a 64-bit architecture and 32 bits if you're on a 32-bit architecture.

Integer literals

- A number literal is a type suffix, such as `57u8`, to designate the type
- Number literals can also use `_` as a visual separator to make the number easier to read, such as `1_000`, which will have the same value as if you had specified `1000`

Number literals	Example
Decimal	<code>let x = 98_222</code>
Hex	<code>let x = 0xff</code>
Octal	<code>let x = 0o77</code>
Binary	<code>let x = 0b1111_0000</code>
Byte (u8 only)	<code>let x = b'A'</code>

Integer Overflow

- Let's say you have a variable of type `u8` that can hold values between 0 and 255.
- If you try to change the variable to a value outside of that range, such as 256, integer overflow will occur, which can result in one of two behaviors:
 - When you're compiling in **debug mode**, Rust includes checks for integer overflow that cause your program to panic at runtime if this behavior occurs.
 - When you're compiling in **release mode** with the `-release` flag, Rust does not include checks for integer overflow that cause panics (In the case of a `u8`, the value 256 becomes 0, the value 257 becomes 1, and so on).

Floating-Point Types

- Rust's floating-point types are `f32` (32 bits) and `f64` (64 bits), see (IEEE-754 STANDARD).
- The default type is `f64` because on modern CPUs it's roughly the same speed as `f32` but is capable of more precision.
- All floating-point types are signed.

```
1 fn main() {  
2     let x = 2.0; // f64  
3     let y: f32 = 3.0; // f32  
4 }
```

```
1 fn main() {  
2     let x = 0.1;  
3     let y = 0.2;  
4     let z = x + y;  
5     println!("result: {0}", z);  
6     ↪ // 0.30000000000000004  
7 }
```

Numeric Operations

```
1 fn main() {  
2     // addition  
3     let sum = 5 + 10;  
4  
5     // subtraction  
6     let difference = 95.5 - 4.3;  
7  
8     // multiplication  
9     let product = 4 * 30;  
10  
11    // division  
12    let quotient = 56.7 / 32.2;  
13    let floored = 2 / 3; // Results in 0  
14  
15    // remainder  
16    let remainder = 43 % 5;  
17 }
```

Boolean Types

Booleans are one byte in size.

```
1 fn main() {  
2     let t = true;  
3  
4     let f: bool = false; // with explicit type annotation  
5 }
```

Character Types

- Note that we specify char literals with **single quotes**, as opposed to string literals, which use double quotes.
- **Rust's char type is four bytes in size** and represents a Unicode Scalar Value, which means it can represent a lot more than just ASCII.

```
fn main() {  
    let c = 'z';  
    let z: char = 'Z'; // with explicit type annotation  
    let heart_eyed_cat = '😺';  
}
```

Tuple Types

A tuple is a general way of grouping together a number of values with a variety of types into one **compound type**.

```
1 fn main() {  
2     let tup = (500, 6.4, 1);  
3     let (x, y, z) = tup;  
4     println!("The value of y is: {y}"); // 6.4  
5  
6     let five_hundred = tup.0;  
7     let six_point_four = tup.1;  
8     let one = tup.2;  
9 }
```

Array Types

- Another way to have a collection of multiple values is with an array.
- Unlike a tuple, every element of **an array must have the same type**.
- Unlike arrays in some other languages, **arrays in Rust have a fixed length**.
- **Rust panics if index out of bounds**.

Note

Arrays are useful when you want your data allocated on the **stack** rather than the **heap**

Array Types (2)

```
1 fn main() {  
2     let a = [1, 2, 3, 4, 5];  
3     let months = ["January", "February", "March", "April", "May",  
4         ↪ "June", "July", "August", "September", "October", "November",  
5         ↪ "December"];  
6     let b: [i32; 5] = [1, 2, 3, 4, 5]; // Here, i32 is the type of  
7         ↪ each element. After the semicolon, the number 5 indicates the  
8         ↪ array contains five elements  
9     let first = a[0];  
10    let second = a[1];  
11 }
```

Functions

Rust code uses **snake case** as the conventional style for function and variable names, in which all letters are lowercase and underscores separate words.

```
1 fn main() {  
2     another_function(5);  
3 }  
4  
5 fn another_function(x: i32) {  
6     println!("The value of x is: {x}");  
7 }
```

functions **parameters** are special variables that are part of a function's signature. When a function has parameters, you can provide it with concrete values for those parameters. Technically, the concrete values are called **arguments**, but in casual conversation, people tend to use the words parameter and argument interchangeably.

Statements vs Expressions

Statements

are instructions that perform some action and do not return a value.

Expressions

evaluate to a resulting value.

```
1 fn main() {  
2     let y = {  
3         let x = 3;  
4         x + 1  
5     };  
6  
7     println!("The value of y is: {y}"); // 4  
8 }
```

Functions with Return Values

- Function bodies are made up of a series of statements optionally ending in an expression.
- Calling a function is an expression.
- Functions can return values to the code that calls them.
- We don't name return values, but we must declare their type after an arrow (`->`).
- In Rust, the **return value of the function** is synonymous with **the value of the final expression in the block of the body** of a function.
- You can return early from a function by using the return keyword and specifying a value, but most functions return the last expression implicitly.

Functions with Return Values (2)

```
1 fn five() -> i32 {  
2     5 // no semicolon because it's an expression whose value we want  
   ↪ to return  
3 }  
4  
5 fn main() {  
6     let x = five();  
7  
8     println!("The value of x is: {x}"); // 6  
9 }
```

Functions with Return Values (3)

```
1 fn main() {  
2     let x = plus_one(5);  
3  
4     println!("The value of x is: {x}");  
5 }  
6  
7 fn plus_one(x: i32) -> i32 {  
8     x + 1; // remove this semicolon  
9 }
```

The definition of the function `plus_one` says that it will return an `i32`, but statements don't evaluate to a value, which is expressed by `()`, the unit type.

Functions with Return Values (4)

```
$ cargo run
   Compiling functions v0.1.0 (file:///projects/functions)
error[E0308]: mismatched types
  --> src/main.rs:7:24
   |
7 | fn plus_one(x: i32) -> i32 {
   |     -----          ^^^ expected `i32`, found `()`
   |     |
   |     implicitly returns `()` as its body has no tail or `return`
   ↳ expression
8 |     x + 1;
   |         - help: remove this semicolon
```

For more information about this error, try `rustc --explain E0308`.
 error: could not compile `functions` due to previous error

if Expressions

if expressions are sometimes called arms

```
1 fn main() {  
2     let number = 6;  
3  
4     if number % 4 == 0 {  
5         println!("number is divisible by 4");  
6     } else if number % 3 == 0 {  
7         println!("number is divisible by 3");  
8     } else if number % 2 == 0 {  
9         println!("number is divisible by 2");  
10    } else {  
11        println!("number is not divisible by 4, 3, or 2");  
12    }  
13 }
```

if Expressions (2)

It's also worth noting that the condition in this code must be a bool. If the condition isn't a bool, we'll get an error.

```

1 fn main() {
2     let number = 3;
3
4     if number {
5         println!("number was
6             ↳ three");
7     }
8 }

```

```

1 fn main() {
2     let number = 3;
3
4     if number != 0 {
5         println!("number was
6             ↳ something other than
7             ↳ zero");
8     }
9 }

```

```
$ cargo run
```

```
...
```

```

4 |         if number {
   |           ~~~~~~ expected `bool`, found integer

```

Using if in a let Statement

Remember that blocks of code evaluate to the last expression in them, and numbers by themselves are also expressions

```
1 fn main() {  
2     let condition = true;  
3     let number = if condition { 5 } else { 6 };  
4  
5     println!("The value of number is: {number}"); // 6  
6 }
```

The results of both the if arm and the else arm should be in the same type. So, `let number = if condition { 5 } else { "six" };` get panic error.

Repetition with Loops

The `loop` keyword tells Rust to execute a block of code over and over again forever or until you explicitly tell it to stop.

```
1 fn main() {  
2     loop {  
3         println!("again!");  
4     }  
5 }
```

Returning Values from Loops

```
1 fn main() {  
2     let mut counter = 0;  
3  
4     let result = loop {  
5         counter += 1;  
6  
7         if counter == 10 {  
8             break counter * 2;  
9         }  
10    };  
11  
12    println!("The result is {result}"); // 20  
13 }
```

Loop Labels to Disambiguate Between Multiple Loops

Loop labels must begin with a single quote

```
1 fn main() {  
2     let mut count = 0;  
3     'counting_up: loop {  
4         println!("count = {count}");  
5         let mut remaining = 10;  
6         loop {  
7             println!("remaining = {remaining}");  
8             if remaining == 9 { break;}  
9             if count == 2 { break 'counting_up; }  
10            remaining -= 1;  
11        }  
12        count += 1;  
13    }  
14    println!("End count = {count}"); // End count = 2  
15 }
```

Conditional Loops with while

```
1 fn main() {  
2     let mut number = 3;  
3  
4     while number != 0 {  
5         println!("{number}!");  
6  
7         number -= 1;  
8     }  
9  
10    println!("LIFTOFF!!!");  
11 }
```

Repetition with for

```
1 fn main() {  
2     let a = [10, 20, 30, 40, 50];  
3  
4     for element in a {  
5         println!("the value is: {element}");  
6     }  
7 }
```

for with range

```
1 fn main() {  
2     for number in (1..4).rev() {  
3         println!("{number}!");  
4     }  
5     println!("LIFTOFF!!!");  
6 }  
7 // output:  
8 // 3!  
9 // 2!  
10 // 1!  
11 // LIFTOFF!!!
```

Understanding Ownership

What Is Ownership?

- **Ownership** is a set of rules that governs how a Rust program manages memory.
- Some languages have **garbage collection** that **regularly looks** (performance?) for no-longer used memory as the program runs;
- In other languages, the **programmer** must **explicitly allocate and free** the memory.
- Rust uses a third approach: memory is managed through a system of ownership with **a set of rules** that the **compiler checks**. If any of the rules are violated, the program won't compile.

Note

None of the features of ownership will slow down your program while it's running.

Because ownership is a new concept for many programmers, it does take some time to get used to.

The Stack and the Heap

- In a systems programming language like Rust, whether a value is on the stack or the heap affects how the language behaves and why you have to make certain decisions. Parts of ownership will be described in relation to the stack and the heap later in the following, so here is a brief explanation.
- Both the **stack** and the **heap** are **parts of memory** available to your code to use at runtime, but they are structured in different ways.
- The **stack** stores values in the order it gets them and removes the values in the opposite order. This is **referred to as last in, first out**. Adding data is called pushing onto the stack, and removing data is called popping off the stack. **All data stored on the stack must have a known, fixed size**. Data with an unknown size at compile time or a size that might change must be stored on the heap instead.

The Stack and the Heap (2)

- The heap is less organized: when you put data on the heap, you request a certain amount of space. The **memory allocator** finds an empty spot in the heap that is big enough, marks it as being in use, and returns a pointer, which is the address of that location. This process is called allocating on the heap and is sometimes abbreviated as just allocating (pushing values onto the stack is not considered allocating). **Because the pointer to the heap is a known, fixed size, you can store the pointer on the stack, but when you want the actual data, you must follow the pointer.**
- **Pushing to the stack is faster than allocating on the heap** because the allocator never has to search for a place to store new data; that location is always at the top of the stack. Comparatively, allocating space on the heap requires more work, because the allocator must first find a big enough space to hold the data and then perform bookkeeping to prepare for the next allocation.

The Stack and the Heap (3)

- When your code calls a function, the values passed into the function (including, potentially, pointers to data on the heap) and the function's local variables get pushed onto the stack. When the function is over, those values get popped off the stack.
- Keeping track of what parts of code are using what data on the heap, minimizing the amount of duplicate data on the heap, and cleaning up unused data on the heap so you don't run out of space are all problems that ownership addresses. Once you understand ownership, you won't need to think about the stack and the heap very often, but knowing that the main purpose of ownership is to manage heap data can help explain why it works the way it does.

Ownership Rules

- Each value in Rust has an owner.
- There can only be one owner at a time.
- When the owner goes out of scope, the value will be dropped.

Variable Scope

```
1 fn main() {  
2     {  
3         // s is not valid here, it's not yet declared  
4         let s = "hello"; // s is valid from this point forward  
5  
6         // do stuff with s  
7     } // this scope is now over, and s is no longer valid  
8 }
```

The String Type

- To illustrate the rules of ownership, we need a data type that is more complex than those we covered in the “Data Types” section
- The types covered previously are all a **known size**, can be **stored on the stack and popped off the stack when their scope is over**, and can be quickly and trivially copied to make a new, independent instance if another part of code needs to use the same value in a different scope.
- We want to look at **data that is stored on the heap** and explore how Rust knows when to clean up that data, and the **String** type is a great example

The String Type (2)

- String literals(`&str`) are hard-coded into our program (**fast and efficient**). String literals are convenient, but they aren't suitable for every situation in which we may want to use text.
 - One reason is that they're **immutable**.
 - Another is that not every string value can be known when we write our code (for example, user input)
- For these situations, Rust has a second string type, `String`.
 - This type manages data allocated on the **heap** and as such is able to store an amount of text that is unknown to us at compile time.
 - You can create a `String` from a string literal using the `from` function, like so:

```
let s = String::from("hello");
```

The String Type (3)

```
1 fn main() {  
2     // (all the type annotations are superfluous)  
3     // A reference to a string allocated in read only memory  
4     let pangram: &'static str = "the quick brown fox jumps over the  
    ↪ lazy dog";  
5     println!("Pangram: {}", pangram);  
6  
7     let mut string = String::new(); // Create an empty and growable  
    ↪ `String`  
8     string.push_str("yumcoder");  
9  
10    // Heap allocate a string  
11    let alice = String::from("I like dogs");  
12    // Allocate new memory and store the modified string there  
13    let bob: String = alice.replace("dog", "cat");  
14    println!("Alice says: {}", alice);  
15    println!("Bob says: {}", bob);  
16 }
```


Move

```
1 fn main() {  
2     let x = 5;  
3     let y = x;  
4     func(x); // y = x, func parameter bind on stack  
5 }  
6 fn func(y: u8) {  
7     y  
8 }
```

Bind the value 5 to x; then make a copy of the value in x and bind it to y.” Because **integers are simple values with a known, fixed size**, and these two 5 values are pushed onto the **stack**

Primitive type on stack

Primitive type’s size known and fixed so pushed onto the **stack**

Move (2)

```

1 let s1 = String::from("hello");
2 let s2 = s1;
3
4 println!("{}", world!", s1); // Error!

```




Figure: a String in memory

The second line would make a copy of the value in `s1` and bind it to `s2`. But this isn't quite what happens.

A String is made up of three parts, shown on the left: a pointer to the memory that holds the contents of the string, a length, and a capacity. This group of data is stored on the stack. On the right is the memory on the heap that holds the contents.

Move (3)

When we assign `s1` to `s2`, the String data is copied, meaning we copy the pointer, the length, and the capacity that are on the stack. We do not copy the data on the heap that the pointer refers to. In other words, the data representation in memory looks like the following (sounds like making a **shallow copy**):





Figure: Representation in memory of the variable `s2` that has a copy of the pointer, length, and capacity of `s1`

Another possibility for what $s2 = s1$ might do if Rust copied the heap data as well

The representation does not look like the following Figure, which is what memory would look like if Rust instead copied the heap data as well. If Rust did this, the operation $s2 = s1$ could be very expensive in terms of runtime performance if the data on the heap were large (**deep copy**).



s1 was moved into s2

- When a variable goes out of scope, **Rust automatically calls the drop function and cleans up the heap memory** for that variable.
- When s2 and s1 go out of scope, they will both try to free the same memory. This is known as a **double free error** and is one of the memory safety bugs.
- **Freeing memory twice** can lead to memory corruption, which can potentially lead to security vulnerabilities.
- To ensure memory safety, after the line `let s2 = s1;`, Rust considers s1 as no longer valid. Therefore, Rust doesn't need to free anything when s1 goes out of scope.

s1 was moved into s2 (2)


```
1 let s1 = String::from("hello");  
2 let s2 = s1;  
3  
4 println!("{}", world!", s1); // Error!
```

s1 was moved into s2 (3)

```
$ cargo run
...
error[E0382]: borrow of moved value: `s1`
  --> src/main.rs:5:28
   |
2  |     let s1 = String::from("hello");
   |           -- move occurs because `s1` has type `String`, which does
   |           ↪ not implement the `Copy` trait
3  |     let s2 = s1;
   |           -- value moved here
4  |
5  |     println!("{}", world!", s1);
   |                               ^^ value borrowed here after move
   |
...
```

s1 was moved into s2 (4)

If you've heard the terms shallow copy and deep copy while working with other languages, the concept of copying the pointer, length, and capacity without copying the data probably sounds like making a shallow copy. But **because Rust also invalidates the first variable, instead of being called a shallow copy, it's known as a move**. In this example, we would say that s1 was moved into s2.



Move notes

Memory safety bugs

That solves our problem! With only `s2` valid, when it goes out of scope it alone will free the memory, and we're done.

Design choice

Rust will never automatically create “deep” copies of your data. Therefore, any automatic copying can be assumed to be inexpensive in terms of runtime performance.

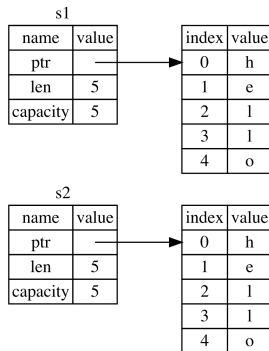
Deeply copy with clone

If we do want to deeply copy the heap data of the String, not just the stack data, we can use a common method called clone.

```

1 let s1 = String::from("hello");
2 let s2 = s1.clone();
3
4 println!("s1 = {}, s2 = {}", s1, s2);

```



Stack-Only Data: Copy and Clone

We don't have a call to clone, but `x` is still valid and wasn't moved into `y`.

```
1 let x = 5;  
2 let y = x;  
3  
4 println!("x = {}, y = {}", x, y);
```

The reason is that types such as integers that have a **known size at compile time** are stored entirely on the **stack**, so **copies of the actual values are quick to make**. That means there's no reason we would want to prevent `x` from being valid after we create the variable `y`. In other words, there's no difference between deep and shallow copying here, so calling `clone` wouldn't do anything different from the usual shallow copying, and we can leave it out.

Ownership and Functions

```

1 fn main() {
2     let s = String::from("hello"); // s comes into scope
3     takes_ownership(s); // s's value moves into the function...
4                           // ... and so is no longer valid here
5     let x = 5; // x comes into scope
6     makes_copy(x); // x would move into the function,
7                     // but i32 is Copy, so it's okay to still use x afterward
8 } // Here, x goes out of scope, then s. But because s's value was moved,
   ↳ nothing
9     // special happens.
10 fn takes_ownership(some_string: String) {
11     // some_string comes into scope
12     println!("{}", some_string);
13 } // Here, some_string goes out of scope and `drop` is called. The backing
   ↳ memory is freed.
14 fn makes_copy(some_integer: i32) {
15     // some_integer comes into scope
16     println!("{}", some_integer);
17 } // Here, some_integer goes out of scope. Nothing special happens.

```

Return Values and Scope

```

1 fn main() {
2     let s1 = gives_ownership(); // gives_ownership moves its return
3                                 // value into s1
4     let s2 = String::from("hello"); // s2 comes into scope
5     let s3 = takes_and_gives_back(s2); // s2 is moved into
6                                         // takes_and_gives_back, which also
7                                         // moves its return value into s3
8 } // Here, s3 goes out of scope and is dropped. s2 was moved, so nothing
9    // happens. s1 goes out of scope and is dropped.
10 fn gives_ownership() -> String {
11     // gives_ownership will move its return value into the function that calls
12     ↪ it
13     let some_string = String::from("yours"); // some_string comes into scope
14     some_string // some_string is returned and moves out to the calling
15     ↪ function
16 }
17 // This function takes a String and returns one
18 fn takes_and_gives_back(a_string: String) -> String {
19     // a_string comes into scope
20     a_string // a_string is returned and moves out to the calling function
21 }

```

Returning ownership

```
1 fn main() {  
2     let s1 = String::from("hello");  
3  
4     let (s2, len) = calculate_length(s1);  
5  
6     println!("The length of '{}' is {}.", s2, len);  
7 }  
8  
9 fn calculate_length(s: String) -> (String, usize) {  
10     let length = s.len(); // len() returns the length of a String  
11  
12     (s, length)  
13 }
```

Rust has a feature for using a value without transferring ownership, called **references**.

References

A reference is **like a pointer** in that it's an **address we can follow to access the data** stored at that address; **that data is owned by some other variable**. Unlike a pointer, **a reference is guaranteed to point to a valid value of a particular type for the life of that reference**.

```

1 fn main() {
2     let s1 = String::from("hello");
3     let len = calculate_length(&s1);
4     println!("The length of '{}' is {}.", s1, len);
5 }
6
7 fn calculate_length(s: &String) -> usize {
8     // s is a reference to a String
9     s.len()
10 } // Here, s goes out of scope. But because it does
    ↪ not have ownership of what it refers to, it is not
    ↪ dropped.

```


Note that we pass `&s1` into `calculate_length` and, in its definition, we take `&String` rather than `String`. These ampersands represent references, and they allow you to refer to some value without taking ownership of it

References (2)

```

1 fn main() {
2     let s1 = String::from("hello");
3     let len = calculate_length(&s1);
4     println!("The length of '{}' is {}.", s1, len);
5 }
6
7 fn calculate_length(s: &String) -> usize {
8     // s is a reference to a String
9     s.len()
10 } // Here, s goes out of scope. But because it does not have ownership of what
    ↳ it refers to, it is not dropped.

```



Borrowing

borrowing: *the action of creating a reference.* As in real life, if a person owns something, you can borrow it from them. When you're done, you have to give it back. You don't own it. So, what happens if we try to modify something we're borrowing?

```
1 fn main() {  
2     let s = String::from("hello");  
3  
4     change(&s);  
5 }  
6  
7 fn change(some_string: &String) {  
8     some_string.push_str(", world"); // Attempting to modify a  
9     ↪ borrowed value!  
10 }
```

Just as variables are immutable by default, so are references. We're not allowed to modify something we have a reference to.

Borrowing (2)

```
$ cargo run
   Compiling ownership v0.1.0 (file:///projects/ownership)
error[E0596]: cannot borrow `*some_string` as mutable, as it is behind
↳ a `&` reference
   --> src/main.rs:8:5
   |
7 | fn change(some_string: &String) {
   |               ----- help: consider changing this to be a
   ↳ mutable reference: `&mut String`
8 |     some_string.push_str(", world");
   |     ~~~~~ `some_string` is a `&`
   ↳ reference, so the data it refers to cannot be borrowed as mutable
```

For more information about this error, try `rustc --explain E0596`.
 error: could not compile `ownership` due to previous error

Borrowing, Mutable References

We can fix the code to allow us to modify a borrowed value with just a few small tweaks that use, instead, a mutable reference (`&mut`).

```
1 fn main() {  
2     let mut s = String::from("hello");  
3  
4     change(&mut s);  
5 }  
6  
7 fn change(some_string: &mut String) {  
8     some_string.push_str(", world");  
9 }
```

Mutable references restriction

If you have a mutable reference to a value, you can have no other references to that value.

```
1 fn main() {  
2     let mut s = String::from("hello");  
3  
4     let r1 = &mut s;  
5     let r2 = &mut s;  
6  
7     println!("{}", r1, r2);  
8 }
```

Mutable references restriction (2)

```
$ cargo run
  Compiling ownership v0.1.0 (file:///projects/ownership)
error[E0499]: cannot borrow `s` as mutable more than once at a time
--> src/main.rs:5:14
|
4 |     let r1 = &mut s;
|               ----- first mutable borrow occurs here
5 |     let r2 = &mut s;
|               ~~~~~ second mutable borrow occurs here
6 |
7 |     println!("{}", r1, r2);
|                   -- first borrow later used here
```

For more information about this error, try ``rustc --explain E0499``.
 error: could not compile `ownership` due to previous error

This error says that this code is invalid because we cannot borrow `s` as mutable more than once at a time. The first mutable borrow is in `r1` and must last until it's used in the `println!`, but between the creation of that mutable reference and its usage, we tried to create another mutable reference in `r2` that borrows the same data as `r1`.

Why mutable references restriction exist in Rust?

It's something that new Rustaceans struggle with because most languages let you mutate whenever you'd like. The benefit of having this restriction is that **Rust can prevent data races at compile time**.

A data race is similar to a race condition and happens when these three behaviors occur:

- Two or more pointers access the same data at the same time.
- At least one of the pointers is being used to write to the data.
- There's no mechanism being used to synchronize access to the data.

Data races cause undefined behavior and can be difficult to diagnose and fix when you're trying to track them down at runtime.

Mutable references restriction example

As always, we can use curly brackets to create a new scope, allowing for multiple mutable references, just not simultaneous ones:

```
1 fn main() {  
2     let mut s = String::from("hello");  
3  
4     {  
5         let r1 = &mut s;  
6     } // r1 goes out of scope here, so we can make a new reference  
       ↪ with no problems.  
7  
8     let r2 = &mut s;  
9 }
```

Mutable references restriction example (combining mutable and immutable references)

```
1 fn main() {  
2     let mut s = String::from("hello");  
3  
4     let r1 = &s; // no problem  
5     let r2 = &s; // no problem  
6     let r3 = &mut s; // BIG PROBLEM  
7  
8     println!("{}", r1, r2, r3);  
9 }
```


Mutable references restriction example (combining mutable and immutable references)

```
$ cargo run
  Compiling ownership v0.1.0 (file:///projects/ownership)
error[E0502]: cannot borrow `s` as mutable because it is also borrowed
↳ as immutable
   --> src/main.rs:6:14
   |
4  |     let r1 = &s; // no problem
   |               -- immutable borrow occurs here
5  |     let r2 = &s; // no problem
6  |     let r3 = &mut s; // BIG PROBLEM
   |               ~~~~~ mutable borrow occurs here
7  |
8  |     println!("{}", r1, r2, r3);
   |                               -- immutable borrow later used here
```

For more information about this error, try `rustc --explain E0502`.
 error: could not compile `ownership` due to previous error



Mutable references restriction example (combining mutable and immutable references, note to reference's scope)

```
1 fn main() {  
2     let mut s = String::from("hello");  
3  
4     let r1 = &s; // no problem  
5     let r2 = &s; // no problem  
6     println!("{}", r1, r2);  
7     // variables r1 and r2 will not be used after this point  
8  
9     let r3 = &mut s; // no problem  
10    println!("{}", r3);  
11 }
```

Dangling References

In languages with pointers, it's easy to erroneously create a dangling pointer—a **pointer that references a location in memory that may have been given to someone else**—by freeing some memory while preserving a pointer to that memory. In Rust, by contrast, **the compiler guarantees that references will never be dangling references**: if you have a reference to some data, the compiler will ensure that the data will not go out of scope before the reference to the data does.

```

1 fn main() {
2     let reference_to_nothing = dangle();
3 }
4
5 fn dangle() -> &String {
6     // dangle returns a reference to a String
7     let s = String::from("hello"); // s is a new String
8     &s // we return a reference to the String, s
9 } // Here, s goes out of scope, and is dropped. Its memory goes away. Danger!

```

Dangling References (2)

```
$ cargo run
  Compiling ownership v0.1.0 (file:///projects/ownership)
error[E0106]: missing lifetime specifier
  --> src/main.rs:5:16
   |
5 | fn dangle() -> &String {
   |               ^ expected named lifetime parameter

= help: this function's return type contains a borrowed value, but
       ↪ there is no value for it to be borrowed from
help: consider using the `static` lifetime
   |
5 | fn dangle() -> &static String {
   |               ~~~~~
```

For more information about this error, try `rustc --explain E0106`.
 error: could not compile `ownership` due to previous error

Dangling References (3), solution

The solution here is to return the `String` directly:

```
1 fn main() {  
2     let string = no_dangle();  
3 }  
4  
5 fn no_dangle() -> String {  
6     let s = String::from("hello");  
7  
8     s  
9 }
```

This works without any problems. **Ownership is moved out, and nothing is deallocated.**

The Slice Type

Slices let you **reference a contiguous sequence of elements in a collection** rather than the whole collection. A slice is **a kind of reference, so it does not have ownership**.

Learn application of Slice Type by example

Here's a small programming problem: write a function that takes a string of words separated by spaces and returns the first word it finds in that string.

```

1 fn first_word(s: &String) -> usize {
2     let bytes = s.as_bytes();
3     for (i, &item) in bytes.iter().enumerate() {
4         if item == b' ' {
5             return i;
6         }
7     }
8     s.len()
9 }
10 fn main() {
11     let mut s = String::from("hello world");
12     let word = first_word(&s); // word will get the value 5
13     s.clear(); // this empties the String, making it equal to ""
14     // word still has the value 5 here, but there's no more string that
15     // we could meaningfully use the value 5 with. word is now totally
16     ↪ invalid!
17 }

```

Learn application of Slice Type by example


- This program compiles without any errors and would also do so if we used `word` after calling `s.clear()`. Because `word` isn't connected to the state of `s` at all, `word` still contains the value 5. We could use that value 5 with the variable `s` to try to extract the first word out, but this would be a bug because the contents of `s` have changed since we saved 5 in `word`.
- Luckily, Rust has a solution to this problem: string slices.

String Slices

```

1 fn main() {
2     let s = String::from("hello world");
3
4     let hello = &s[0..5];
5     let world = &s[6..11];
6
7     let slice = &s[0..2];
8     let slice = &s[..2];
9
10    let slice = &s[3..len];
11    let slice = &s[3..];
12 }

```



Rewrite first_word to return a slice

```
1 fn first_word(s: &String) -> &str {
2     let bytes = s.as_bytes();
3
4     for (i, &item) in bytes.iter().enumerate() {
5         if item == b' ' {
6             return &s[0..i];
7         }
8     }
9
10    &s[..]
11 }
12
13 fn main() {
14     let mut s = String::from("hello world");
15     let word = first_word(&s);
16     s.clear(); // error!
17     println!("the first word is: {}", word);
18 }
```

Rewrite first_word to return a slice (2)

```
$ cargo run
  Compiling ownership v0.1.0 (file:///projects/ownership)
error[E0502]: cannot borrow `s` as mutable because it is also borrowed
↳ as immutable
   --> src/main.rs:18:5
   |
16 |     let word = first_word(&s);
   |                               -- immutable borrow occurs here
17 |
18 |     s.clear(); // error!
   |     ^^^^^^^^^ mutable borrow occurs here
19 |
20 |     println!("the first word is: {}", word);
   |                                           ---- immutable borrow later
   ↳ used here
```

For more information about this error, try `rustc --explain E0502`.
 error: could not compile `ownership` due to previous error

Rewrite `first_word` to return a slice (3)

Recall from the borrowing rules that if we have an immutable reference to something, we cannot also take a mutable reference. Because `clear` needs to truncate the `String`, it needs to get a mutable reference. The `println!` after the call to `clear` uses the reference in `word`, so the immutable reference must still be active at that point. Rust disallows the mutable reference in `clear` and the immutable reference in `word` from existing at the same time, and compilation fails. Not only has Rust made our API easier to use, but it has also eliminated an entire class of errors at compile time!

Improve first_word function

```
fn first_word(s: &String) -> &str { /*...*/ }
```

A more experienced Rustacean would write the signature as in below instead because it allows us to use the same function on both `&String` values and `&str` values.

```
fn first_word(s: &str) -> &str { /*...*/ }
```

Using Structs to Structure Related Data

Defining and Instantiating Structs

- A **struct, or structure**, is a custom data type that lets you package together and name multiple related values that make up a meaningful group. If you're familiar with an object-oriented language, a struct is like an *object's data attributes*.
- **Structs are similar to tuples**, in that both hold multiple related values. Like tuples, the pieces of a struct can be different types. Unlike with tuples, in a struct you'll name each piece of data so it's clear what the values mean. Adding these names means that structs are more flexible than tuples: you don't have to rely on the order of the data to specify or access the values of an instance.

Defining and Instantiating Structs (2)

```
1 struct User {
2     active: bool,
3     username: String,
4     email: String,
5     sign_in_count: u64,
6 }
7 fn main() {
8     let mut user1 = User {
9         email: String::from("someone@example.com"),
10        username: String::from("someusername123"),
11        active: true,
12        sign_in_count: 1,
13    };
14
15    user1.email = String::from("anotheremail@example.com");
16 }
```


Using the Field Init Shorthand

```
1 fn build_user(email: String, username: String) -> User {
2     User {
3         email: email,
4         username: username,
5         active: true,
6         sign_in_count: 1,
7     }
8 }
9 fn build_user_shorthand(email: String, username: String) -> User {
10     User {
11         email,
12         username,
13         active: true,
14         sign_in_count: 1,
15     }
16 }
```

Creating Instances From Other Instances With Struct Update Syntax

```
1 fn main() {
2     // --snip--
3     let user2 = User {
4         active: user1.active,
5         username: user1.username,
6         email: String::from("another@example.com"),
7         sign_in_count: user1.sign_in_count,
8     };
9 }
10 fn main_less_code() {
11     // --snip--
12     let user2 = User {
13         email: String::from("another@example.com"),
14         ..user1
15     };
16 }
```

Creating Instances From Other Instances With Struct Update Syntax (2)

The `..user1` **must come last to specify that any remaining fields should get their values from the corresponding fields** in `user1`.

Note that the struct update syntax uses `=` like an assignment

This is because it moves the data. In this example, we can no longer use `user1` after creating `user2` because the `String` in the `username` field of `user1` was moved into `user2`. If we had given `user2` new `String` values for both `email` and `username`, and thus only used the `active` and `sign_in_count` values from `user1`, then `user1` would still be valid after creating `user2`. The types of `active` and `sign_in_count` are types that implement the `Copy` trait.

Creating Instances From Other Instances With Struct Update Syntax (3)

```
1 let mut user1 = User {
2     email: String::from("someone@example.com"),
3     username: String::from("someusername123"),
4     active: true,
5     sign_in_count: 1,
6 };
7
8
9 let user2 = User {
10     email: String::from("another@example.com"),
11     ..user1
12 };
13 println!("{:?}", user1);
```

Creating Instances From Other Instances With Struct Update Syntax (4)

```

error[E0382]: borrow of partially moved value: `user1`
  --> src/main.rs:22:23
   |
18 |         let user2 = User {
   |         -----^
19 |             email: String::from("another@example.com"),
20 |             ..user1
21 |         };
   |         |_____- value partially moved here
22 |         println!("{::#?}", user1);
   |                                ~~~~~ value borrowed here after partial
   |                                move
   |
= note: partial move occurs because `user1.username` has type
   |     `String`, which does not implement the `Copy` trait

```

Ownership of Struct Data

- In the User struct definition, **we used the owned String type rather than the `&str` string slice type**. This is a deliberate choice because we want each instance of this struct to own all of its data and for that data to be valid for as long as the entire struct is valid.
- **It's also possible for structs to store references to data owned by something else, but to do so requires the use of lifetimes**, a Rust feature that we'll discuss later. Lifetimes ensure that the data referenced by a struct is valid for as long as the struct is.

Using Tuple Structs without Named Fields to Create Different Types

Rust also supports structs that look similar to tuples, called tuple structs. Tuple structs have the added meaning the struct name provides but don't have names associated with their fields; rather, they just have the types of the fields. Tuple structs are useful when you want to give the whole tuple a name and make the tuple a different type from other tuples, and when naming each field as in a regular struct would be verbose or redundant.

```
1 struct Color(i32, i32, i32);
2 struct Point(i32, i32, i32);
3
4 fn main() {
5     let black = Color(0, 0, 0);
6     let origin = Point(0, 0, 0);
7 }
```

Unit-Like Structs Without Any Fields

You can also define structs that don't have any fields! These are called **unit-like structs** because they behave similarly to ().

```
1 struct AlwaysEqual;  
2  
3 fn main() {  
4     let subject = AlwaysEqual;  
5 }
```

Unit type

The tuple without any values has a special name, `unit`. This value and its corresponding type are both written `()` and represent an empty value or an empty return type. Expressions implicitly return the unit value if they don't return any other value.

Unit-Like Structs Without Any Fields (2)

Unit-like structs can be useful when you need to implement a trait on some type but don't have any data that you want to store in the type itself. Example, see: The global memory allocator, `Global`, is a unit struct.

Method Syntax

Methods are similar to **functions**: we declare them with the `fn` keyword and a name, they can have parameters and a return value, and they contain some code that's run when the method is called from somewhere else. Unlike functions, methods are defined within the context of a struct (or an enum or a trait object, which we cover later), and their **first parameter is always self**, which represents the instance of the struct the method is being called on.

Method Syntax(2)

```
1  #[derive(Debug)]
2  struct Rectangle {
3      width: u32,
4      height: u32,
5  }
6  impl Rectangle {
7      fn area(&self) -> u32 {
8          self.width * self.height
9      }
10 }
11 fn main() {
12     let rect1 = Rectangle {
13         width: 30,
14         height: 50,
15     };
16     println!("The area of the rectangle is {} square pixels.",
17         ↵ rect1.area());
17 }
```

Associated Functions

- **All functions defined within an `impl` block** are called associated functions because they're associated with the type named after the `impl`. We can define associated functions that don't have self as their first parameter (and thus are not methods) because they don't need an instance of the type to work with. We've already used one function like this: the `String::from` function that's defined on the `String` type.
- Associated functions that aren't methods are often used for constructors that will return **a new instance of the struct**(factory method). These are often called **new**, but `new` isn't a special name and isn't built into the language.

Associated Functions

```
1 #[derive(Debug)]
2 struct Rectangle {
3     width: u32,
4     height: u32,
5 }
6
7 impl Rectangle {
8     fn square(size: u32) -> Self {
9         Self {
10             width: size,
11             height: size,
12         }
13     }
14 }
15 fn main() {
16     let sq = Rectangle::square(3);
17 }
```

Multiple `impl` Blocks

```
1 impl Rectangle {  
2     fn area(&self) -> u32 {  
3         self.width * self.height  
4     }  
5 }  
6  
7 impl Rectangle {  
8     fn can_hold(&self, other: &Rectangle) -> bool {  
9         self.width > other.width && self.height > other.height  
10    }  
11 }
```

Enums and Pattern Matching

Defining an Enum

Enums allow you to define a type by enumerating its possible variants.

Struct vs Enum

Where structs give you a way of grouping together related fields and data, like a Rectangle with its width and height, enums give you a way of saying a value is one of a possible set of values.

```
1 enum IpAddrKind {
2     V4,
3     V6,
4 }
5 let four = IpAddrKind::V4;
6 let six = IpAddrKind::V6;
7
8 fn route(ip_kind: IpAddrKind) {}
9 route(IpAddrKind::V4);
10 route(IpAddrKind::V6);
```


Defining an Enum (2)

There's another advantage to using an enum rather than a struct: **each variant can have different types and amounts of associated data.**

```
1 enum IpAddr {  
2     V4(u8, u8, u8, u8),  
3     V6(String),  
4 }  
5  
6 let home = IpAddr::V4(127, 0, 0, 1);  
7  
8 let loopback = IpAddr::V6(String::from("::1"));
```

Defining an Enum (3)

```
1 enum Message {  
2     Quit, // has no data associated with it at all.  
3     Move { x: i32, y: i32 }, // has named fields like a struct does.  
4     Write(String), // includes a single String.  
5     ChangeColor(i32, i32, i32), // includes three i32 values.  
6 }  
7  
8 // defining different kinds of struct  
9 struct QuitMessage; // unit struct  
10 struct MoveMessage {  
11     x: i32,  
12     y: i32,  
13 }  
14 struct WriteMessage(String); // tuple struct  
15 struct ChangeColorMessage(i32, i32, i32); // tuple struct
```

The Option Enum and Its Advantages Over Null Values

The Option type encodes the very common scenario in which a value could be something or it could be nothing.

```
1 enum Option<T> {  
2     None,  
3     Some(T),  
4 }  
5  
6 let some_number = Some(5);  
7 let some_char = Some('e');  
8  
9 let absent_number: Option<i32> = None;
```

see: “Null References: The Billion Dollar Mistake”

Advantage of Option Over Null

- When we have a `Some` value, we know that a value is present and the value is held within the `Some`. When we have a `None` value, in some sense, it means the same thing as null: we don't have a valid value. So why is having `Option<T>` any better than having null?
 - In short, because `Option<T>` and `T` (where `T` can be any type) are different types, the compiler won't let us use an `Option<T>` value as if it were definitely a valid value. For example, this code won't compile because it's trying to add an `i8` to an `Option<i8>`:

```

1 let x: i8 = 5;
2 let y: Option<i8> = Some(5);
3
4 let sum = x + y; // Error!
5 // no implementation for `i8 + Option<i8>`
6 // the trait `Add<Option<i8>>` is not implemented for `i8`

```

The match Control Flow Construct

Rust has an extremely powerful control flow construct called `match` that allows you to compare a value against a series of patterns and then execute code based on which pattern matches.

```
1 enum Coin {  
2     Penny,  
3     Nickel,  
4     Dime,  
5     Quarter,  
6 }  
7  
8 fn value_in_cents(coin: Coin) -> u8 {  
9     match coin {  
10         Coin::Penny => 1,  
11         Coin::Nickel => 5,  
12         Coin::Dime => 10,  
13         Coin::Quarter => 25,  
14     }  
15 }
```



Patterns that Bind to Values

```

1  enum UsState {
2      Alabama,
3      Alaska,
4      // --snip--
5  }
6  enum Coin {
7      Penny,
8      Nickel,
9      Dime,
10     Quarter(UsState),
11 }
12 fn value_in_cents(coin: Coin) -> u8 {
13     match coin {
14         Coin::Penny => 1,
15         Coin::Nickel => 5,
16         Coin::Dime => 10,
17         Coin::Quarter(state) => {
18             println!("State quarter from {:?}!", state);
19             25
20         }
21     }
22 }

```

Another useful feature of match arms is that they can bind to the parts of the values that match the pattern. This is how we can extract values out of enum variants.

Matching with `Option<T>`

```
1 fn plus_one(x: Option<i32>) -> Option<i32> {  
2     match x {  
3         None => None,  
4         Some(i) => Some(i + 1),  
5     }  
6 }  
7  
8 let five = Some(5);  
9 let six = plus_one(five);  
10 let none = plus_one(None);
```

Matches Are Exhaustive

Matches in Rust are exhaustive: we must exhaust every last possibility in order for the code to be valid. Especially in the case of `Option<T>`, when Rust prevents us from forgetting to explicitly handle the `None` case, it protects us from assuming that we have a value when we might have null, thus making the billion-dollar mistake impossible.

```

1 fn plus_one(x: Option<i32>) -> Option<i32> {
2     match x {
3         Some(i) => Some(i + 1),
4     }
5 }
6 // Error! We didn't handle the None case, so this code will cause a
  ↪ bug. Rust knows that we didn't cover every possible case and even
  ↪ knows which pattern we forgot!
7 // 3    ~    Some(i) => Some(i + 1),
8 // 4    ~    None => todo!(),

```


Catch-all Patterns and the `_` Placeholder

```
1 let dice_roll = 9;
2 match dice_roll {
3     3 => add_fancy_hat(),
4     7 => remove_fancy_hat(),
5     other => move_player(other),
6 }
7
8 match dice_roll {
9     3 => add_fancy_hat(),
10    7 => remove_fancy_hat(),
11    _ => reroll(),
12 }
```

Concise Control Flow with if let

The if let syntax lets you combine if and let into a less verbose way to handle values that match one pattern while ignoring the rest.

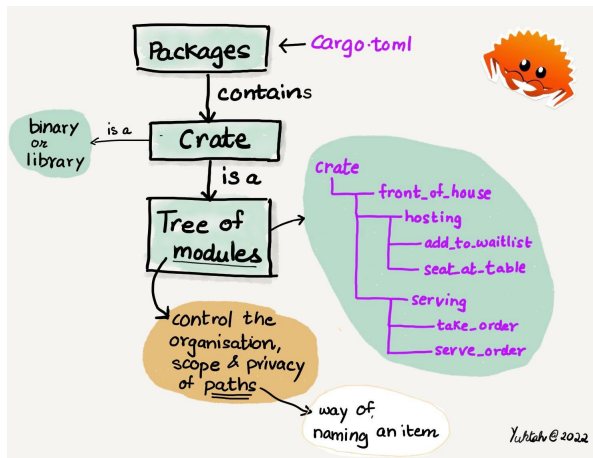
```
1 let config_max = Some(3u8);
2 match config_max {
3     Some(max) => println!("The maximum is configured to be {}", max),
4     _ => (),
5 }
6
7 if let Some(max) = config_max {
8     println!("The maximum is configured to be {}", max);
9 }
```

Managing Projects with Packages, Crates, and Modules

Packages and Crates

- As a project grows, you should organize code by splitting it into multiple **modules** and then multiple files.
- A **package** can contain multiple binary crates and optionally one library crate.
- As a package grows, you can extract parts into separate crates that become external dependencies.
- A **crate** is the smallest amount of code that the Rust compiler considers at a time. A crate can come in one of two forms:
 - **Binary crates** are programs you can compile to an executable, such as a command-line program or a server. Each must have a function called `main` that defines what happens when the executable runs.
 - **Library crates** don't have a main function, and they don't compile to an executable. Instead, they define functionality intended to be shared with multiple projects.
- For very large projects comprising a set of interrelated packages that evolve together, Cargo provides **workspaces**, which we'll cover later.

Module system



Package

- A package is a bundle of one or more crates that provides a set of functionality. A package contains a `Cargo.toml` file that describes how to build those crates.
- After we run `cargo new`, we use `ls` to see what Cargo creates. In the project directory, there's a `Cargo.toml` file, giving us a package. There's also a `src` directory that contains `main.rs`. Cargo knows that if the package directory contains `src/lib.rs`, the package contains a library crate with the same name as the package, and `src/lib.rs` is its crate root.
- If a package contains `src/main.rs` and `src/lib.rs`, it has two crates: a binary and a library, both with the same name as the package.

Declaring modules

In the crate root file, you can **declare new modules**; say, you declare a “garden” module with `mod garden;`. The compiler will look for the module’s code in these places:

- **Inline**, within curly brackets that replace the semicolon following `mod garden`
- **In the file** `src/garden.rs`
- **In the file** `src/garden/mod.rs`

Declaring submodules: In any file other than the crate root, you can declare submodules. For example, you might declare `mod vegetables;` in `src/garden.rs`. The compiler will look for the submodule’s code within the directory named for the parent module in these places:

- **Inline**, directly following `mod vegetables`, within curly brackets instead of the semicolon
- **In the file** `src/garden/vegetables.rs`
- **In the file** `src/garden/vegetables/mod.rs`

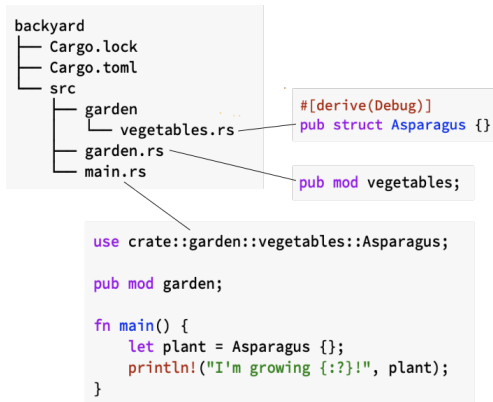
Control the privacy with modules

- Modules let us organize code within a crate for readability and easy reuse.
- Modules also **allow us to control the privacy of items**, because **code within a module is private by default**. Private items are internal implementation details not available for outside use.

Paths to code in modules

Once a module is part of your crate, you can refer to code in that module from anywhere else in that same crate, as long as the privacy rules allow, using the path to the code. For example, an `Asparagus` type in the `garden` `vegetables` module would be found at

```
crate::garden::vegetables::Asparagus .
```



Paths in the Module

A path can take two forms:

- An **absolute** path is the full path starting from a crate root; for code from an external crate, **the absolute path begins with the crate name**, and for code from the current crate, it starts with the literal `crate`.
- A **relative** path starts from the current module and **uses `self`, `super`, or an identifier in the current module**.

Both absolute and relative paths are followed by one or more identifiers separated by double colons (`::`).

Paths in the Module (2)

```
1 mod front_of_house {
2     pub mod hosting {
3         pub fn add_to_waitlist() {} // public
4         fn private_add_to_waitlist() {} // private
5     }
6 }
7
8 pub fn eat_at_restaurant() {
9     // Absolute path
10    crate::front_of_house::hosting::add_to_waitlist();
11
12    // Relative path
13    front_of_house::hosting::add_to_waitlist();
14 }
```

Both `front_of_house` module and `add_to_waitlist` method are public. The `pub` keyword lets us use these paths in `add_to_waitlist` with respect to the privacy rules.

Siblings are friend

While `front_of_house` isn't public, because the `eat_at_restaurant` function is defined in the same module as `front_of_house` (that is, `eat_at_restaurant` and `front_of_house` are siblings), we can refer to `front_of_house` from `eat_at_restaurant`.

Starting Relative Paths with `super`

```
1 fn deliver_order() {}  
2  
3 mod back_of_house {  
4     fn fix_incorrect_order() {  
5         cook_order();  
6         super::deliver_order();  
7     }  
8  
9     fn cook_order() {}  
10 }
```

Making Structs Public

```

1 mod back_of_house {
2     pub struct Breakfast {
3         pub toast: String,
4         seasonal_fruit: String, // private
5     }
6     impl Breakfast {
7         pub fn summer(toast: &str) -> Breakfast {
8             Breakfast {
9                 toast: String::from(toast),
10                seasonal_fruit: String::from("peaches"),
11            }
12        }
13    }
14 }
15 pub fn eat_at_restaurant() {
16     // Order a breakfast in the summer with Rye toast
17     let mut meal = back_of_house::Breakfast::summer("Rye");
18     // Change our mind about what bread we'd like
19     meal.toast = String::from("Wheat");
20     println!("I'd like {} toast please", meal.toast);
21     // The next line won't compile if we uncomment it; we're not allowed
22     // to see or modify the seasonal fruit that comes with the meal
23     // meal.seasonal_fruit = String::from("blueberries");
24 }

```

Making Enums Public

```
1 mod back_of_house {  
2     pub enum Appetizer {  
3         Soup,  
4         Salad,  
5     }  
6 }  
7  
8 pub fn eat_at_restaurant() {  
9     let order1 = back_of_house::Appetizer::Soup;  
10    let order2 = back_of_house::Appetizer::Salad;  
11 }
```

Bringing Paths into Scope with the `use` Keyword

- Having to write out the paths to call functions can feel inconvenient and repetitive.
- Fortunately, there's a way to simplify this process: we can create a shortcut to a path with the `use` keyword once, and then `use` the shorter name everywhere else in the **scope**.

```
1 mod front_of_house {  
2     pub mod hosting {  
3         pub fn add_to_waitlist() {}  
4     }  
5 }  
6 use crate::front_of_house::hosting; // scope: root module  
7 pub fn eat_at_restaurant() {  
8     hosting::add_to_waitlist();  
9 }
```

Note that `use` only creates the shortcut for the particular scope in which the use occurs.

Bringing Paths into Scope with the `use` Keyword (2)

```

1 mod front_of_house {
2     pub mod hosting {
3         pub fn add_to_waitlist() {}
4     }
5 }
6
7 use crate::front_of_house::hosting;
8
9 mod customer {
10     pub fn eat_at_restaurant() {
11         hosting::add_to_waitlist(); // Error!
12         // ~~~~~ use of undeclared crate or module `hosting`
13     }
14 }

```

To fix this problem, move the `use` within the `customer` module too, or reference the shortcut in the parent module with `super::hosting` within the child `customer` module.

Creating Idiomatic `use` Paths

```

1 mod front_of_house {
2     pub mod hosting {
3         pub fn add_to_waitlist() {}
4     }
5 }
6 use crate::front_of_house::hosting; // idiomatic way!
7 pub fn eat_at_restaurant() {
8     hosting::add_to_waitlist();
9 }

```

```

1 mod front_of_house {
2     pub mod hosting {
3         pub fn add_to_waitlist() {}
4     }
5 }
6 use crate::front_of_house::hosting::add_to_waitlist;
7 pub fn eat_at_restaurant() {
8     add_to_waitlist();
9 }

```

Creating Idiomatic `use` Paths (2)

- Bringing the function's parent module into scope with `use` means we have to specify the parent module when calling the function.

Specifying the parent module when calling the function makes it clear that the function isn't locally defined while still minimizing repetition of the full path.

- On the other hand, when bringing in structs, enums, and other items with `use`, it's idiomatic to specify the full path. There's no strong reason behind this idiom: it's just the convention that has emerged, and folks have gotten used to reading and writing Rust code this way.

```
1 use std::collections::HashMap;
2
3 fn main() {
4     let mut map = HashMap::new();
5     map.insert(1, 2);
6 }
```

Creating Idiomatic `use` Paths (3)

```
1 use std::fmt;
2 use std::io;
3 fn function1() -> fmt::Result {
4     // --snip--
5 }
6 fn function2() -> io::Result<()> {
7     // --snip--
8 }
```

Idiomatic way

- **Function:** bringing the **function's parent** module into scope.
- **Structs, enums, and other items:** it's idiomatic to specify the **full path**.
 - The exception to this idiom is if we're bringing two items with the same name into scope with `use` statements, because Rust doesn't allow that.

Providing New Names with the as Keyword

```
1 use std::fmt::Result;
2 use std::io::Result as IoResult;
3
4 fn function1() -> Result {
5     // --snip--
6 }
7
8 fn function2() -> IoResult<()> {
9     // --snip--
10 }
```

Re-exporting Names with `pub use`

When we bring a name into scope with the `use` keyword, the name available in the new scope is private. To enable the code that calls our code to refer to that name as if it had been defined in that code's scope, we can combine `pub` and `use`.

```

1 mod front_of_house {
2     pub mod hosting {
3         pub fn add_to_waitlist() {}
4     }
5 }
6
7 pub use crate::front_of_house::hosting;
8
9 pub fn eat_at_restaurant() {
10     hosting::add_to_waitlist();
11 }

```

It makes our library well organized for programmers working on the library and programmers calling the library.

Using External Packages

Members of the Rust community have made many packages available at `crates.io`, and pulling any of them into your package involves these same steps: listing them in your package's `Cargo.toml` file and using `use` to bring items from their crates into scope. Example:

1) Add the following line to `Cargo.toml` file.

```
rand = "0.8.5"
```

2) bring items from the crates into scope (with `use`) and used it.

```
1 use rand::Rng;
2
3 fn main() {
4     let secret_number = rand::thread_rng().gen_range(1..=100);
5 }
```

Using Nested Paths to Clean Up Large `use` Lists

```
1 use std::cmp::Ordering;
2 use std::io;
3 // we can use nested paths to bring the same items into scope in one
  ↳ line.
4 use std::{cmp::Ordering, io};
5
6 use std::io;
7 use std::io::Write;
8 // To merge these two paths into one use statement, we can use self in
  ↳ the nested path
9 use std::io::{self, Write};
10
11 // If we want to bring all public items defined in a path into scope,
  ↳ we can specify that path followed by the * glob operator
12 use std::collections::*;
```


Separating Modules into Different Files

- For a module named `front_of_house` declared in the crate root, the compiler will look for the module's code in:
 - `src/front_of_house.rs` (what we covered)
 - `src/front_of_house/mod.rs` (older style, still supported path)
- For a module named `hosting` that is a submodule of `front_of_house`, the compiler will look for the module's code in:
 - `src/front_of_house/hosting.rs` (what we covered)
 - `src/front_of_house/hosting/mod.rs` (older style, still supported path)
- If you use both styles for the same module, you'll get a compiler error. Using a mix of both styles for different modules in the same project is allowed, but might be confusing for people navigating your project.
- The main downside to the style that uses files named `mod.rs` is that your project can end up with many files named `mod.rs`, which is confusing.

Common Collections

Collections

- **Rust's standard library** includes a number of very useful **data structures** called **collections**.
- Most other data types represent one specific value, but collections can contain multiple values.
- Unlike the built-in array and tuple types, the data these collections point to is stored on the **heap**, which means the amount of data does not need to be known at compile time and can grow or shrink as the program runs.
- Each kind of collection has different capabilities and costs, and choosing an appropriate one for your current situation is a skill you'll develop over time.
- Rust's collections can be grouped into four major categories:
 - **Sequences**: Vec, VecDeque(double-ended queue), LinkedList
 - **Maps**: HashMap, BTreeMap
 - **Sets**: HashSet, BTreeSet
 - **Misc**: BinaryHeap

Creating a New Vector

```
1 fn main() {  
2     let v1: Vec<i32> = Vec::new(); // We added a type annotation here,  
    ↪ because we aren't inserting any values into this vector, Rust  
    ↪ doesn't know what kind of elements we intend to store  
3  
4     // More often, you'll create a Vec<T> with initial values and Rust  
    ↪ will infer the type of value you want to store, so you rarely  
    ↪ need to do this type annotation. Rust conveniently provides  
    ↪ the vec! macro, which will create a new vector that holds the  
    ↪ values you give it.  
5     let v2 = vec![1, 2, 3];  
6  
7     let mut v3 = Vec::new();  
8     v3.push(5);  
9     v3.push(6);  
10 }
```

Reading Elements of Vectors

There are two ways to reference a value stored in a vector: via indexing or using the `get` method.

```
1 fn main() {  
2     let v = vec![1, 2, 3, 4, 5];  
3  
4     let third: &i32 = &v[2];  
5     println!("The third element is {third}");  
6  
7     let third: Option<&i32> = v.get(2);  
8     match third {  
9         Some(third) => println!("The third element is {third}"),  
10        None => println!("There is no third element."),  
11    }  
12 }
```

The reason Rust provides these two ways to reference an element is so you can choose how the program behaves when you try to use an index value outside the range of existing elements.

Iterating over the Values in a Vector

```
1 fn main() {  
2     let v = vec![100, 32, 57];  
3     for i in &v {  
4         println!("{i}");  
5     }  
6  
7     for i in &mut v {  
8         *i += 50; // we have to use the * dereference operator to get  
9         ↳ to the value in i before we can use the += operator  
10    }
```

Using an Enum to Store Multiple Types

```
1 fn main() {  
2     enum SpreadsheetCell {  
3         Int(i32),  
4         Float(f64),  
5         Text(String),  
6     }  
7  
8     let row = vec![  
9         SpreadsheetCell::Int(3),  
10        SpreadsheetCell::Text(String::from("blue")),  
11        SpreadsheetCell::Float(10.12),  
12    ];  
13 }
```

Using an Enum to Store Multiple Types

```
1 fn main() {  
2     enum SpreadsheetCell {  
3         Int(i32),  
4         Float(f64),  
5         Text(String),  
6     }  
7  
8     let row = vec![  
9         SpreadsheetCell::Int(3),  
10        SpreadsheetCell::Text(String::from("blue")),  
11        SpreadsheetCell::Float(10.12),  
12    ];  
13 }
```


What Is a `String` ?

- Many of the same operations available with `Vec<T>` are available with `String` as well, because `String` is actually implemented as a wrapper around **a vector of bytes** with some extra guarantees, restrictions, and capabilities.
- Rust has only one **string type** *in the core language*, which is the **string slice** `str` that is usually seen in its borrowed form `&str`.
- The **`String` type**, which is provided by **Rust's standard library** rather than coded into the core language, is *a growable, mutable, owned, UTF-8 encoded string type*. When Rustaceans refer to “strings” in Rust, they might be referring to either the `String` or the string slice `&str` types, not just one of those types.

Creating a New String

```
1 fn main() {  
2     let mut s0 = String::new();  
3  
4     let data = "initial contents";  
5     let s1 = data.to_string();  
6  
7     // the method also works on a literal directly:  
8     let s2 = "initial contents".to_string();  
9  
10    let s3 = String::from("initial contents");  
11 }
```

Because strings are used for so many things, we can use many different generic APIs for strings, providing us with a lot of options. Some of them can seem redundant, but they all have their place! In this case, `String::from` and `to_string` do the same thing, so which you choose is a matter of style and readability.

Updating a String

```
1 fn main() {  
2     let mut s0 = String::new();  
3  
4     let data = "initial contents";  
5     let s1 = data.to_string();  
6  
7     // the method also works on a literal directly:  
8     let s2 = "initial contents".to_string();  
9  
10    let s3 = String::from("initial contents");  
11 }
```

Hash Maps

```
1 use std::collections::HashMap;
2
3 let mut scores = HashMap::new();
4
5 scores.insert(String::from("Blue"), 10);
6 scores.insert(String::from("Yellow"), 50);
7
8 let team_name = String::from("Blue");
9 let score = scores.get(&team_name).copied().unwrap_or(0);
10 // This program handles the Option by calling copied to get an
11    ↳ Option<i32> rather than an Option<&i32>, then unwrap_or to set
12    ↳ score to zero if scores doesn't have an entry for the key.
13
14 for (key, value) in &scores {
15     println!("{key}: {value}");
16 }
```

Hash Maps and Ownership

For types that implement the Copy trait, like i32, the values are copied into the hash map. For owned values like String, the values will be moved and the hash map will be the owner of those values.

```
1 use std::collections::HashMap;
2
3 let field_name = String::from("Favorite color");
4 let field_value = String::from("Blue");
5
6 let mut map = HashMap::new();
7 map.insert(field_name, field_value);
8 // field_name and field_value are invalid at this point, try using
   ↳ them and see what compiler error you get!
```

Error Handling

Rust errors categories

Rust groups errors into two major categories: **recoverable** and **unrecoverable** errors.

- For a **recoverable** error, such as a file not found error, we most likely just want to report the problem to the user and retry the operation.
- **Unrecoverable** errors are always **symptoms of bugs**, like trying to access a location beyond the end of an array, and so we want to immediately **stop the program**.
- Rust doesn't have exceptions. Instead, it has the type `Result<T, E>` for recoverable errors and the `panic!` macro that stops execution when the program encounters an unrecoverable error.

Unrecoverable Errors with panic!

```
1 fn main() {  
2     let v = vec![1, 2, 3];  
3     v[99]; // panic  
4  
5     panic!("crash and burn"); // panic macro  
6 }
```

Unwinding the Stack

By default, when a panic occurs, the program starts unwinding, which means Rust walks back up the stack and cleans up the data from each function it encounters. However, this walking back and cleanup is a lot of work. Rust, therefore, allows you to choose the alternative of immediately aborting, which ends the program without cleaning up.

Recoverable Errors with Result

Most errors aren't serious enough to require the program to stop entirely. Sometimes, when a function fails, it's for a reason that you can easily interpret and respond to (For example, open a file that doesn't exist).

```
1 enum Result<T, E> {  
2     Ok(T),  
3     Err(E),  
4 }
```

```
1 use std::fs::File;  
2  
3 fn main() {  
4     let greeting_file_result = File::open("hello.txt");  
5  
6     let greeting_file = match greeting_file_result {  
7         Ok(file) => file,  
8         Err(error) => panic!("Problem opening the file: {:?}", error),  
9     };  
10 }
```

Matching on Different Errors

```
1 use std::fs::File;
2 use std::io::ErrorKind;
3
4 fn main() {
5     let greeting_file_result = File::open("hello.txt");
6
7     let greeting_file = match greeting_file_result {
8         Ok(file) => file,
9         Err(error) => match error.kind() {
10             ErrorKind::NotFound => match File::create("hello.txt") {
11                 Ok(fc) => fc,
12                 Err(e) => panic!("Problem creating the file: {:?}", e),
13             },
14             other_error => {
15                 panic!("Problem opening the file: {:?}", other_error);
16             }
17         },
18     };
19 }
```

Alternatives to Using match with `Result<T, E>`

```
1 use std::fs::File;
2 use std::io::ErrorKind;
3
4 fn main() {
5     let greeting_file = File::open("hello.txt").unwrap_or_else(|error| {
6         if error.kind() == ErrorKind::NotFound {
7             File::create("hello.txt").unwrap_or_else(|error| {
8                 panic!("Problem creating the file: {:?}", error);
9             })
10         } else {
11             panic!("Problem opening the file: {:?}", error);
12         }
13     });
14 }
```

Shortcuts for Panic on Error: `unwrap` and `expect`

The `Result<T, E>` type has many helper methods defined on it to do various, more specific tasks. If the `Result` value is the `Ok` variant, `unwrap` will return the value inside the `Ok`. If the `Result` is the `Err` variant, `unwrap` will call the `panic!` macro for us.

```
1 use std::fs::File;
2
3 fn main() {
4     let greeting_file = File::open("hello.txt").unwrap();
5 }
```

We use `expect` in the same way as `unwrap`: to return the file handle or call the `panic!` macro.

```
1 use std::fs::File;
2
3 fn main() {
4     let greeting_file =
5         File::open("hello.txt").expect("hello.txt should be included in this
6         ↪ project");
7 }
```

Propagating Errors with ? Operator

```
1 use std::fs::File;
2 use std::io::{self, Read};
3
4 fn read_username_from_file() -> Result<String, io::Error> {
5     let username_file_result = File::open("hello.txt");
6
7     let mut username_file = match username_file_result {
8         Ok(file) => file,
9         Err(e) => return Err(e),
10    };
11
12    let mut username = String::new();
13
14    match username_file.read_to_string(&mut username) {
15        Ok(_) => Ok(username),
16        Err(e) => Err(e),
17    }
18 }
```

A Shortcut for Propagating Errors: the `?` Operator

```

1 use std::fs::File;
2 use std::io::{self, Read};
3
4 fn read_username_from_file() -> Result<String, io::Error> {
5     let mut username_file = File::open("hello.txt")?; // If the value
6     ↳ of the Result is an Ok, the value inside the Ok will get
7     ↳ returned from this expression, and the program will continue.
8     ↳ If the value is an Err, the Err will be returned from the
9     ↳ whole function as if we had used the return keyword so the
10    ↳ error value gets propagated to the calling code
11    let mut username = String::new();
12    username_file.read_to_string(&mut username)?;
13
14    Ok(username)
15 }

```

A Shortcut for Propagating Errors: the `?` Operator

```

1 use std::fs::File;
2 use std::io::{self, Read};
3
4 fn read_username_from_file() -> Result<String, io::Error> {
5     let mut username = String::new();
6
7     File::open("hello.txt")?.read_to_string(&mut username)?;
8
9     Ok(username)
10 }
```

```

1 use std::fs;
2 use std::io;
3
4 fn read_username_from_file() -> Result<String, io::Error> {
5     // Reading a file into a string is a fairly common operation, so the
6     ↳ standard library provides the convenient fs::read_to_string function.
7     fs::read_to_string("hello.txt")
8 }
```

Where The `?` Operator Can Be Used

- The `?` operator can only be used in functions whose return type is compatible with the value the `?` is used on.
- we're only allowed to use the `?` operator in a function that returns `Result`, `Option`, or another type that implements `FromResidual`.

```

1 use std::fs::File;
2
3 fn main() {
4     let greeting_file = File::open("hello.txt");
5     // Error!
6     // ^ cannot use the `?` operator in a function that returns `()`
7     // help: the trait `FromResidual<Result<Infallible,
8     //      ↪ std::io::Error>>` is not implemented for `()`
9 }

```


Where The `?` Operator Can Be Used (2)

```
1 use std::error::Error;  
2 use std::fs::File;  
3  
4 fn main() -> Result<(), Box<dyn Error>> {  
5     let greeting_file = File::open("hello.txt");  
6  
7     Ok(())  
8 }
```

The `Box<dyn Error>` type is a trait object, which we'll talk about later. For now, you can read `Box<dyn Error>` to mean "any kind of error."

Generic Types, Traits, and Lifetimes

Generic Data Types

We use generics to create definitions for items like function signatures or structs, which we can then use with many different concrete data types.

```

1 fn largest_i32(list: &[i32]) ->
  ↳ &i32 {
2     let mut largest = &list[0];
3
4     for item in list {
5         if item > largest {
6             largest = item;
7         }
8     }
9
10    largest
11 }
  
```

```

1 fn largest_char(list: &[char]) ->
  ↳ &char {
2     let mut largest = &list[0];
3
4     for item in list {
5         if item > largest {
6             largest = item;
7         }
8     }
9
10    largest
11 }
  
```

Generic Data Types (3)

```
1 fn largest<T>(list: &[amp;T]) -> &T {
2     let mut largest = &list[0];
3
4     for item in list {
5         if item > largest {
6             largest = item;
7         }
8     }
9
10    largest
11 }
12
13 fn main() {
14     let number_list = vec![34, 50, 25, 100, 65];
15     let result = largest(&number_list);
16     println!("The largest number is {}", result);
17
18     let char_list = vec!['y', 'm', 'a', 'q'];
19     let result = largest(&char_list);
20     println!("The largest char is {}", result);
21 }
```

Generic Data Types In Struct Definitions

```

1 struct Point<T> {
2     x: T,
3     y: T,
4 }
5
6 fn main() {
7     let integer = Point { x: 5, y: 10 };
8     let float = Point { x: 1.0, y: 4.0 };
9
10    // The fields x and y must be the same type because both have the
11    ↪ same generic data type T
12    // let wont_work = Point { x: 5, y: 4.0 }; // expected integer,
13    ↪ found floating-point number
14    // struct Point<T, U> {
15    //     x: T,
16    //     y: U,
17    // }
18 }

```

Generic Data Types In Enum Definitions

```
1 enum Option<T> {  
2     Some(T),  
3     None,  
4 }  
5  
6 enum Result<T, E> {  
7     Ok(T),  
8     Err(E),  
9 }
```

Generic Data Types In Method Definitions

```

1 struct Point<T> {
2     x: T,
3     y: T,
4 }
5 impl<T> Point<T> {
6     fn x(&self) -> &T {
7         &self.x
8     }
9 }
10 // We can also specify constraints on generic types when defining methods on
   ↳ the type. This code means the type Point<f32> will have a
   ↳ distance_from_origin method; other instances of Point<T> where T is not of
   ↳ type f32 will not have this method defined.
11 impl Point<f32> {
12     fn distance_from_origin(&self) -> f32 {
13         (self.x.powi(2) + self.y.powi(2)).sqrt()
14     }
15 }
16 fn main() {
17     let p = Point { x: 5, y: 10 };
18     println!("p.x = {}", p.x());
19 }

```



Generic Data Types In Method Definitions(2)

```
1 struct Point<X1, Y1> {
2     x: X1,
3     y: Y1,
4 }
5 impl<X1, Y1> Point<X1, Y1> {
6     fn mixup<X2, Y2>(self, other: Point<X2, Y2>) -> Point<X1, Y2> {
7         Point {
8             x: self.x,
9             y: other.y,
10        }
11    }
12 }
13 fn main() {
14     let p1 = Point { x: 5, y: 10.4 };
15     let p2 = Point { x: "Hello", y: 'c' };
16
17     let p3 = p1.mixup(p2);
18
19     println!("p3.x = {}, p3.y = {}", p3.x, p3.y);
20 }
```


Monomorphization

Monomorphization is the process of turning generic code into specific code by filling in the concrete types that are used when **compiled**. In this process, the compiler does the opposite of the steps we used to create the generic function: the compiler looks at all the places where generic code is called and generates code for the concrete types the generic code is called with.

Monomorphization

```
1 let integer = Some(5);
2 let float = Some(5.0);
3 // The monomorphized version of the above code looks similar to the following
  ↳ (the compiler uses different names than what we're using here for
  ↳ illustration):
4 enum Option_i32 {
5     Some(i32),
6     None,
7 }
8
9 enum Option_f64 {
10     Some(f64),
11     None,
12 }
13
14 fn main() {
15     let integer = Option_i32::Some(5);
16     let float = Option_f64::Some(5.0);
17 }
```

Traits: Defining Shared Behavior

- A trait defines functionality a particular type has and can share with other types. We can use traits **to define shared behavior in an abstract way**. We can use trait bounds to specify that a generic type can be any type that has certain behavior.
- Note: Traits are similar to a feature often called **interfaces in other languages**, although with some differences.
- Trait definitions are a way to group method signatures together to define a set of behaviors necessary to accomplish some purpose.

Traits: Defining Shared Behavior (2)

```
1 pub trait Summary {  
2     fn summarize(&self) -> String;  
3 }  
4  
5 pub struct Tweet {  
6     pub username: String,  
7     pub content: String,  
8     pub reply: bool,  
9     pub retweet: bool,  
10 }  
11  
12 impl Summary for Tweet {  
13     fn summarize(&self) -> String {  
14         format!("{}: {}", self.username, self.content)  
15     }  
16 }
```

Traits: Default Implementations

```
1 pub trait Summary {
2     fn summarize(&self) -> String {
3         String::from("(Read more...)")
4     }
5 }
6
7 pub struct NewsArticle {
8     pub headline: String,
9     // --snip--
10 }
11 impl Summary for NewsArticle {}
12
13 pub struct Tweet {
14     pub username: String,
15     // --snip--
16 }
17 impl Summary for Tweet {
18     fn summarize(&self) -> String {
19         format!("{}: {}", self.username, self.content)
20     }
21 }
```

Traits as Parameters

```

1 pub fn notify(item: &impl Summary) {
2     println!("Breaking news! {}", item.summarize());
3 }
4 // The impl Trait syntax works for straightforward cases but is
5   ↳ actually syntax sugar for a longer form known as a trait bound; it
6   ↳ looks like this:
7 pub fn notify<T: Summary>(item: &T) {
8     println!("Breaking news! {}", item.summarize());
9 }
10 // Specifying Multiple Trait Bounds with the + Syntax
11 pub fn notify(item: &(impl Summary + Display)) {}
12 // The + syntax is also valid with trait bounds on generic types:
13 pub fn notify<T: Summary + Display>(item: &T) {}

```

Clearer Trait Bounds with where Clauses

A functions with multiple generic type parameters can contain lots of trait bound information between the function's name and its parameter list, making the function signature hard to read. For this reason, Rust has alternate syntax for specifying trait bounds inside a where clause after the function signature. So instead of writing this:

```
1 | fn some_function<T: Display + Clone, U: Clone + Debug>(t: &T, u: &U) -> i32 {}
```

we can use a where clause, like this:

```
1 | fn some_function<T, U>(t: &T, u: &U) -> i32
2 | where
3 |     T: Display + Clone,
4 |     U: Clone + Debug,
5 | {
```

Validating References with Lifetimes

- lifetimes **ensure that references are valid as long as we need them to be.**
- Lifetimes are another **kind of generic** that we've already been using.
- Most of the time, lifetimes are implicit and inferred, just like most of the time, types are inferred. We only must annotate types when multiple types are possible.
- Annotating lifetimes is not even a concept most other programming languages have, so this is going to feel **unfamiliar**.
- The Rust compiler has a **borrow checker** that compares scopes to determine whether all borrows are valid.
 - borrowing: the action of creating a reference.

Preventing Dangling References with Lifetimes

Annotations of the lifetimes of `r` and `x`, named `'a` and `'b`, respectively

```

1 fn main() {
2     let r;                                // -----+-- 'a
3                                           //          |
4     {                                    //          |
5         let x = 5;                       // -+-- 'b |
6         r = &x;                          // |      |
7     }                                    // -+      |
8                                           //          |
9     println!("r: {}", r);               //          |
10 }                                       // -----+
  
```

The variable `x` doesn't "live long enough." The reason is that `x` will be out of scope when the inner scope ends on line 7. But `r` is still valid for the outer scope; because its scope is larger, we say that it "lives longer."

Preventing Dangling References with Lifetimes(2)

```
$ cargo run
  Compiling chapter10 v0.1.0 (file:///projects/chapter10)
error[E0597]: `x` does not live long enough
  --> src/main.rs:6:13
   |
6 |         r = &x;
   |           ^^ borrowed value does not live long enough
7 |     }
   |     - `x` dropped here while still borrowed
8 |
9 |     println!("r: {}", r);
   |                       - borrow later used here
```

For more information about this error, try ``rustc --explain E0597``.
 error: could not compile `chapter10` due to previous error

Generic Lifetimes in Functions

We'll write a function that returns the longer of two string slices. This function will take two string slices and return a single string slice.

```
1 fn main() {  
2     let string1 = String::from("abcd");  
3     let string2 = "xyz";  
4  
5     let result = longest(string1.as_str(), string2);  
6     println!("The longest string is {}", result);  
7 }
```

Generic Lifetimes in Functions (2)

Rust can't tell whether the reference being returned refers to `x` or `y`. Actually, we don't know either, because the `if` block in the body of this function returns a reference to `x` and the `else` block returns a reference to `y`!

```
1 fn longest(x: &str, y: &str) -> &str {  
2     if x.len() > y.len() {  
3         x  
4     } else {  
5         y  
6     }  
7 }
```

Generic Lifetimes in Functions (3)

```
$ cargo run
  Compiling chapter10 v0.1.0 (file:///projects/chapter10)
error[E0106]: missing lifetime specifier
  --> src/main.rs:9:33
   |
9 | fn longest(x: &str, y: &str) -> &str {
   |           ----      ----      ^ expected named lifetime parameter
   |
= help: this function's return type contains a borrowed value, but
       ↪ the signature does not say whether it is borrowed from `x` or `y`
help: consider introducing a named lifetime parameter
   |
9 | fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {
   |           +++++      ++              ++      ++
```

For more information about this error, try ``rustc --explain E0106``.
 error: could not compile `chapter10` due to previous error

Lifetime Annotation Syntax

- Lifetime annotations **don't change how long any of the references live**. Rather, they **describe the relationships of the lifetimes of multiple references to each other** without affecting the lifetimes.
- Lifetime annotations have a slightly **unusual syntax**: the names of lifetime parameters must start with an apostrophe (') and are usually all lowercase and very short, like generic types. Most people use the name `'a` for the first lifetime annotation. We place lifetime parameter annotations after the `&` of a reference, using a space to separate the annotation from the reference's type.

```

1 | &i32           // a reference
2 | &'a i32        // a reference with an explicit lifetime
3 | &'a mut i32    // a mutable reference with an explicit lifetime

```

Lifetime Annotations in Function Signatures

```

1 fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {
2     if x.len() > y.len() {
3         x
4     } else {
5         y
6     }
7 }

```

The generic lifetime `'a` will get the concrete **lifetime that is equal to the smaller of the lifetimes** of `x` and `y`. In practice, it means that the lifetime of the reference returned by the `longest` function is the same as the smaller of the lifetimes of the values referred to by the function arguments. These relationships are what we want Rust to use when analyzing this code.

Lifetime Annotations in Function Signatures (2)

```
1 fn main() {  
2     let string1 = String::from("long string is long");  
3     let result;  
4     {  
5         let string2 = String::from("xyz");  
6         result = longest(string1.as_str(), string2.as_str());  
7     }  
8     println!("The longest string is {}", result);  
9 }
```

The error shows that for `result` to be valid for the `println!` statement, `string2` would need to be valid until the end of the outer scope. Rust knows this because we annotated the lifetimes of the function parameters and return values using the same lifetime parameter `'a`.

Lifetime Annotations in Function Signatures (3)

```
$ cargo run
  Compiling chapter10 v0.1.0 (file:///projects/chapter10)
error[E0597]: `string2` does not live long enough
  --> src/main.rs:6:44
   |
6 |         result = longest(string1.as_str(), string2.as_str());
   |                                     ~~~~~
   |                                     borrowed value does not live long enough
7 |     }
   |     - `string2` dropped here while still borrowed
8 |     println!("The longest string is {}", result);
   |                                     ----- borrow later used
   |     ↪ here
```

For more information about this error, try `rustc --explain E0597`.
 error: could not compile `chapter10` due to previous error

Lifetime Annotations in Struct Definitions

```
1 struct ImportantExcerpt<'a> {  
2     part: &'a str,  
3 }  
4  
5 fn main() {  
6     let novel = String::from("Call me Ishmael. Some years ago...");  
7     let first_sentence = novel.split('.').next().expect("Could not  
8     ↪ find a '.'");  
9     let i = ImportantExcerpt {  
10         part: first_sentence,  
11     };  
12 }
```

We can define structs to hold references, but in that case we would need to add a lifetime annotation on every reference in the struct's definition.

Lifetime Elision

- In early versions (pre-1.0) of Rust, every reference needed an explicit lifetime. After writing a lot of Rust code, the Rust team found that Rust programmers were entering the same lifetime annotations over and over in particular situations. These situations were predictable and followed a few deterministic patterns. After writing a lot of Rust code, the Rust team found that Rust programmers were entering **the same lifetime annotations over and over in particular situations. These situations were predictable and followed a few deterministic patterns.**
- **The patterns programmed into Rust's analysis of references are called the lifetime elision rules.** These aren't rules for programmers to follow; they're a set of particular cases that the compiler will consider, and if your code fits these cases, you don't need to write the lifetimes explicitly.

Lifetime Elision (2)

- The first rule is that the compiler assigns a lifetime parameter to each parameter that's a reference. In other words, a function with one parameter gets one lifetime parameter: `fn foo<'a>(x: &'a i32);` a function with two parameters gets two separate lifetime parameters: `fn foo<'a, 'b>(x: &'a i32, y: &'b i32);` and so on.
- The second rule is that, if there is exactly one input lifetime parameter, that lifetime is assigned to all output lifetime parameters: `fn foo<'a>(x: &'a i32) -> &'a i32 .`
- The third rule is that, if there are multiple input lifetime parameters, but one of them is `&self` or `&mut self` because this is a method, the lifetime of `self` is assigned to all output lifetime parameters. This third rule makes methods much nicer to read and write because fewer symbols are necessary.

Lifetime Elision (2)

```

1 fn first_word(s: &str) -> &str {} // Ok, lifetime elision rules are applied!
2 // Then the compiler applies the first rule, which specifies that each
  ↳ parameter gets its own lifetime. We'll call it 'a as usual, so now the
  ↳ signature is this:
3 fn first_word<'a>(s: &'a str) -> &str {}
4 // The second rule applies because there is exactly one input lifetime. The
  ↳ second rule specifies that the lifetime of the one input parameter gets
  ↳ assigned to the output lifetime, so the signature is now this:
5 fn first_word<'a>(s: &'a str) -> &'a str {}
6 //
7 // example 2:
8 fn longest(x: &str, y: &str) -> &str {}
9 fn longest<'a, 'b>(x: &'a str, y: &'b str) -> &str {}
10 // You can see that the second rule doesn't apply because there is more than
  ↳ one input lifetime. The third rule doesn't apply either, because longest is
  ↳ a function rather than a method, so none of the parameters are self. After
  ↳ working through all three rules, we still haven't figured out what the
  ↳ return type's lifetime is. This is why we got an error trying to compile
  ↳ the code.

```

Lifetime Annotations in Method Definitions

When we implement methods on a struct with lifetimes, we use the same syntax as that of generic type parameters shown:

```

1  impl<'a> ImportantExcerpt<'a> {
2      fn announce_and_return_part(&self, announcement: &str) -> &str {
3          println!("Attention please: {}", announcement);
4          self.part
5      }
6  }

```

The lifetime parameter declaration after `impl` and its use after the type name are required, but we're not required to annotate the lifetime of the reference to `self` because of the first elision rule.

The Static Lifetime

One special lifetime we need to discuss is `'static`, which denotes that the affected reference can live for the entire duration of the program. All string literals have the `'static` lifetime, which we can annotate as follows:

```
1 | let s: &'static str = "I have a static lifetime.";
```

The text of this string is stored directly in the program's binary, which is always available. Therefore, the lifetime of all string literals is `'static`.

Generic Type Parameters, Trait Bounds, and Lifetimes Together

```
1 use std::fmt::Display;
2
3 fn longest_with_an_announcement<'a, T>(x: &'a str, y: &'a str, ann: T)
4   ↪ -> &'a str
5 where
6     T: Display,
7 {
8     println!("Announcement! {}", ann);
9     if x.len() > y.len() {
10         x
11     } else {
12         y
13     }
14 }
```


Writing Automated Tests

The Anatomy of a Test Function

```
1 #[cfg(test)]
2 mod tests {
3     #[test]
4     fn it_works() {
5         let result = 2 + 2;
6         assert_eq!(result, 4);
7     }
8 }
```

```
$ cargo test
```

```
...
```

```
running 1 test
```

```
test tests::it_works ... ok
```

```
test result: ok. 1 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out;
```

```
↳ finished in 0.00s
```

```
...
```

Adding a second test that will fail because we call the `panic!` macro

```
1  #[cfg(test)]
2  mod tests {
3      #[test]
4      fn exploration() {
5          assert_eq!(2 + 2, 4);
6      }
7
8      #[test]
9      fn another() {
10         panic!("Make this test fail");
11     }
12 }
```

Checking Results with the `assert!` Macro

The `assert!` macro, provided by the standard library, is useful when you want to ensure that some condition in a test evaluates to true.

```
1 #[cfg(test)]
2 mod tests {
3     use super::*;
4
5     #[test]
6     fn larger_can_hold_smaller() {
7         let larger = Rectangle {
8             width: 8,
9             height: 7,
10        };
11        let smaller = Rectangle {
12            width: 5,
13            height: 1,
14        };
15
16        assert!(larger.can_hold(&smaller));
17    }
18 }
```

Checking Results with the `assert_eq!` and `assert_ne!` Macro

```
1 pub fn add_two(a: i32) -> i32 {  
2     a + 2  
3 }  
4  
5 #[cfg(test)]  
6 mod tests {  
7     use super::*;  
8  
9     #[test]  
10    fn it_adds_two() {  
11        assert_eq!(4, add_two(2));  
12    }  
13 }
```

Adding Custom Failure Messages

You can also add a custom message to be printed with the failure message as optional arguments to the `assert!`, `assert_eq!`, and `assert_ne!` macros. Any arguments specified after the required arguments are passed along to the `format!` macro

```

1 pub fn greeting(name: &str) -> String {
2     format!("Hello {}!", name)
3 }
4 #[cfg(test)]
5 mod tests {
6     use super::*;
7     #[test]
8     fn greeting_contains_name() {
9         let result = greeting("Carol");
10        assert!(
11            result.contains("Carol"),
12            "Greeting did not contain name, value was `{}`",
13            result
14        );
15    }
16 }

```



Checking for Panics with `should_panic`

```
1 pub struct Guess {
2     value: i32,
3 }
4 impl Guess {
5     pub fn new(value: i32) -> Guess {
6         if value < 1 || value > 100 {
7             panic!("Guess value must be between 1 and 100, got {}. ", value);
8         }
9         Guess { value }
10    }
11 }
12 #[cfg(test)]
13 mod tests {
14     use super::*;
15     #[test]
16     #[should_panic]
17     fn greater_than_100() {
18         Guess::new(200);
19     }
20 }
```

Using `Result<T, E>` in Tests

```
1 #[cfg(test)]
2 mod tests {
3     #[test]
4     fn it_works() -> Result<(), String> {
5         if 2 + 2 == 4 {
6             Ok(())
7         } else {
8             Err(String::from("two plus two does not equal four"))
9         }
10    }
11 }
```

Writing tests so they return a `Result<T, E>` enables you to use the question mark operator in the body of tests, which can be a convenient way to write tests that should fail if any operation within them returns an `Err` variant.

Controlling How Tests Are Run

- `cargo test` compiles your code in test mode and runs the resulting test binary. The default behavior of the binary produced by `cargo test` is to **run all the tests in parallel** and capture output generated during test runs, **preventing the output from being displayed** and making it easier to read the output related to the test results. You can, however, specify command line options to change this default behavior.
- If you don't want to run the tests in parallel or if you want more fine-grained control over the number of threads used, you can send the `--test-threads` flag and the number of threads you want to use to the test binary (ex: `cargo test -- --test-threads=1`).
- If we call `println!` in a test, we won't see the `println!` output in the terminal. If we want to see printed values for passing tests, we can tell Rust to also show the output of successful tests with `cargo test -- --show-output`.

Controlling How Tests Are Run(2)

- Sometimes, running a full test suite can take a long time:
 - If we run the tests without passing any arguments, all the tests will run in parallel
 - We can pass the name of any test function to `cargo test` to run only that test.
 - We can specify part of a test name, and any test whose name matches that value will be run.
- Sometimes a few specific tests can be very time-consuming to execute, so you might want to exclude them during most runs of `cargo test`. After `#[test]` we add the `#[ignore]` line to the test we want to exclude. If we want to run only the ignored tests, we can use `cargo test -- --ignored`.

```
1 #[test]
2 #[ignore]
3 fn expensive_test() {
4     // code that takes an hour to run
5 }
```

Test Organization

- The Rust community thinks about tests in terms of two main categories: **unit tests** and **integration tests**.
 - Unit tests are small and more focused, **testing one module in isolation at a time**, and **can test private interfaces**. You'll put unit tests in the `src` directory in each file with the code that they're testing. The convention is to create a module named `tests` in each file to contain the test functions and to annotate the module with `#[cfg(test)]`.
 - Integration tests are entirely external to your library and use your code in the same way any other external code would, **using only the public interface and potentially exercising multiple modules per test**. Their purpose is to test **whether many parts of your library work together correctly**.

Integration test

We create a tests directory at the top level of our project directory, next to src. Cargo knows to look for integration test files in this directory. We can then make as many test files as we want, and Cargo will compile each of the files as an individual crate.

Filename: `tests/integration_test.rs`

```

addrer
├── Cargo.lock
├── Cargo.toml
├── src
│   └── lib.rs
└── tests
    └── integration_test.rs
  
```

```

1  // Each file in the tests directory is a
   ↳ separate crate, so we need to bring our
   ↳ library into each test crate's scope.
2  use addrer;
3
4  #[test]
5  fn it_adds_two() {
6      assert_eq!(4, addrer::add_two(2));
7  }
  
```

```
$ cargo test
```

```
$ cargo test --test integration_test
```

Submodules in Integration Tests

For example, if we create `tests/common.rs` and place a function named `setup` in it, we can add some code to `tsetup` that we want to call from multiple test functions in multiple test files:

```

├── Cargo.lock
├── Cargo.toml
├── src
│   └── lib.rs
├── tests
│   ├── common
│   │   └── mod.rs
│   └── integration_test.rs

```

Filename: `tests/common.rs`

```

1 pub fn setup() {
2     // setup code specific to your library's
3     ↪ tests would go here
4 }

```

Filename: `tests/integration_test.rs`

```

1 use adder;
2 mod common;
3 #[test]
4 fn it_adds_two() {
5     common::setup();
6     assert_eq!(4, adder::add_two(2));
7 }

```

An I/O Project: Building a Command Line Program

Functional Language Features: Iterators and Closures

More About Cargo and Crates.io

Smart Pointers

Fearless Concurrency

Object-Oriented Programming Features of Rust

Patterns and Matching

Advanced Features

Multithread Web Server

Tokio

Thank you!